

Minutes Analysis

Connor Krenzer

1/4/2021

Introduction

The Federal Reserve is the central bank of the United States. Central banks conduct a country's monetary policy (policy pertaining to money-interest rates, regulation of the banking sector, etc.) in the aims of achieving it's 'dual mandate' of full employment (loosely defined as "anyone who wants a job can get a job") and price stability (predictable price changes throughout the United States as a whole). You can read more about what the Fed is and what they do on their website. For all of you Fed haters who think I'm misleading the public by linking to the Fed's website, do not disparage! Perhaps one of these articles is more your style.

Every six weeks, policy makers on the Federal Reserve's Federal Open Market Committee (FOMC) meet to discuss economic indicators to determine and vote on policy in response to these indicators, among other things. The FOMC Statement released following this meeting can send the entire economy into turmoil—so much so that the minutes are withheld from the public for three weeks before becoming available.

The data in this analysis includes all FOMC Minutes between January 28th, 2009 to November 5th, 2020 (inclusive). As of the time of writing, the Minutes for the December 15th-16th, 2020 meeting have not been released.

I will explore these texts using various text mining approaches including sentiment analysis, tf-idf, and topic modeling to identify trends within the FOMC's Meetings.

Packages

I recommend pulling the information using the HTML version of these meetings. The PDFs are a nightmare to clean (speaking from experience)! If you do decide to read the pdf text into R, I suggest using the `extract_text()` function from the `tabulizer` package because it is smart enough to recognize columns of selectable PDF text. Other packages popular for reading in PDFs, such as `readtext` or `pdftools`, will only read across the page left to right, then down—despite the pages containing two separate columns of text.

Data Import

Since I started this process using the pdf files, I will use their file names to determine the corresponding url. If I had started mining this data from HTML in the first place, a more elegant solution might have been used, but downloading all the PDF files and using their names to put together urls works fine enough. . .

There are two formats for the HTML file—one for the Fed Minutes before 2012 and those released on or after that year. With this in mind, two functions will be used to read in the different formats. The documents are parsed in chronological order but are also arranged after being added to the tibble, in the event that something were to happen to the order of the files in the "Fed Minutes Releases" folder from which these names were pulled.

The Dataset

The data is almost ready for analysis. We now have two columns: one for the document name, and another for the contents of that document. If you are interested in performing a similar analysis, the data frame can now be transformed in a number of ways. You can use the `unnest_tokens()` function from the `tidytext` package to extract tokens (words, groups of words, paragraphs, etc.) and perform sentiment analysis, or you could make a term-document matrix and perform topic modeling.

We will do both, but let's begin by taking a look at the dataset:

```
millions of dollars equivalent)bank of canada bank of mexico any changes in the terms of existing
swap arrangements and the proposed terms of any new arrangements that may be authorized
shall be referred for review and approval to the committee. . all transactions in foreign currencies
undertaken under paragraph .a. abo
```


The quote above is an excerpt from the Minutes release for the meeting dated January 28th, 2009. You might notice that some words have parsing irregularities, such as the lack of space between characters on the parentheses surrounding “millions of dollars equivalent.” This section of text was chosen to represent the ‘worst-case-scenario’ for parsing text. No algorithm is perfect, but one can see these errors are infrequent and therefore unlikely to heavily skew the results of the analysis. The text is from a table detailing relations with the Bank of Canada and the Bank of Mexico—tables typically have issues parsing.

A great place to start any textual analysis is with word counts.

Common Words

Brian F. Madigan is the Secretary who writes all these meetings—he has been the secretary for at least the last 12 years! In a way, we are exploring one man's work over the course of a decade...

After removing ‘stop words’—those words that are extremely common and provide very little meaning (Ex. “the,” “of,” “a,” “to,” etc.)—which words does Brian use most frequently?



These are words one expects to find in official banking documents; the terms “committee,” “inflation,” and “market” dominate across all 95 meetings. They may be insightful to some, but most people would agree these terms are uninteresting finds. Much of what Brian wrote is generic banking lingo from which little meaning can be derived. We can still use these terms to draw interesting conclusions, however; we will visualize their usage over time, for example, but we first want to know which terms are unique to specific documents.

To extract words that have most significance in each document, we can add a weighting factor to terms with the tf-idf statistic.

TF-IDF

The tf-idf statistic is designed to reduce the importance of words that occur often across documents. The math behind it is based on the natural logarithm. It is the product of the ‘term frequency’ and ‘inverse document frequency’:

$$\text{Term Frequency} = (\# \text{ of Occurrences in Document}) / (\# \text{ of Unique Words in Document})$$

$$\text{Inverse Document Frequency} = \ln((\text{Total } \# \text{ of Docs in Corpus}) / (\# \text{ of Docs where Term Appears}))$$

Therefore... **tf-idf = (tf * idf)**


You can think of tf-idf as a formula that finds terms that occur frequently—but not too frequently. Let’s see the most important words among Fed Minutes releases, as measured by the tf-idf:

Table 1: Important Words as Measured by tf-idf

doc_id	word	n	tf	idf	tf_idf
fomcminutes20201105	pandemic	35	0.0037977	2.762117	0.0104898
fomcminutes20200315	coronavirus	34	0.0034729	2.607967	0.0090573
fomcminutes20200916	pandemic	32	0.0032673	2.762117	0.0090247
fomcminutes20200729	pandemic	25	0.0028620	2.762117	0.0079053
fomcminutes20200315	outbreak	29	0.0029622	2.607967	0.0077253
fomcminutes20171101	hurricanes	22	0.0029283	2.607967	0.0076368
fomcminutes20200429	outbreak	24	0.0027574	2.607967	0.0071911
fomcminutes20200610	yct	15	0.0015247	4.553877	0.0069433
fomcminutes20200429	coronavirus	23	0.0026425	2.607967	0.0068915
fomcminutes20170920	hurricanes	18	0.0020829	2.607967	0.0054320

I’m tired of seeing those words everywhere. The correct interpretation of this table, however, is that “pandemic,” “coronavirus,” and “outbreak” are the most important terms, as determined by the tf-idf. Let’s filter out the words “pandemic,” “coronavirus,” “virus,” and “outbreak” to see what else tf-idf deems important. Finally, after some inspection, the names of the months, “shall,” and “our” should also be filtered out.

Let’s see a plot of these terms after the removal:



That’s better. The Fed uses many different acronyms to describe its programs. To name a few, “yct” is the “yield control target,” talb is short for “term asset-backed securities loan facilities,” and “elb” stands for “effective lower bound.” In terms of weather, 2017 was a tough year for the Atlantic and Gulf Coasts of the United States—so much so that “hurricanes” (Hurricanes Harvey, Irma, and Maria formed in 2017—the






FOMC mentioned Harvey and Irma primarily) took two of the top 10 slots. The Fed took measures to provide relief in the aftermath of the storms, which explains the word's prominence.

Aside from the obvious, you may be wondering how these terms relate to the current events at the time. Yield curve controls (yield curve targets, YCT), for instance, were a hot topic in June 2020, when the Fed evaluated the experiences of central banks across the developed world in controlling government bond yields in response to the recession. As a matter of fact, “yct” only appeared in *one* Fed Minutes release and appeared fifteen times! The second phrase, term asset-backed securities loan facilities (talf), was a major policy the Fed pursued in the aftermath of the Great Recession.

Clearly, then, we can see the value brought from this statistic. It allows us to identify important aspects of documents to hone in on topics of interest.

Change Over Time

In addition to finding which terms are important to specific documents (FOMC Meetings in this instance), we can also calculate the usage of words over time. We can find some cool trends, such as the following:




Of course, we don't have to restrict our comparisons of words used in FOMC meetings to other words in FOMC meetings. Why don't we compare the Fed's use of "inflation" with 10 year bond yields?

Note: The words "employment" and "unemployment" have similar usage over time, so I picked "employment" for the graph. Their numbers could be pooled, but the purpose of these graphs is only to introduce ways of finding trends in the dataset.

Note: Only a handful of FOMC meetings started on the first day of the month in the sample period, so the index's base will be set to 100 on November 1st, 2017. Most economic data only comes out monthly, quarterly, etc., and a common base is needed for comparison.

FOMC Mentions of "Inflation" vs. 10 Year Treasury Maturity Rate




As we can see, FOMC mentions of inflation and the constant maturity rate seem to be inversely related.

N-GRAMS


We do not need to restrict our analysis to one word. We can extract multiple words, or "n-grams," to add context to the phrases the Fed uses. We can do the same analysis as in the previous sections with two or three words instead. After some investigation, know that the words "federal," "reserve," "committee," and "intermeeting" are also removed, as I suspect they mostly refer to formalities within the documents and not necessarily to the topics discussed at FOMC meetings.

First, the most common pairs of words:

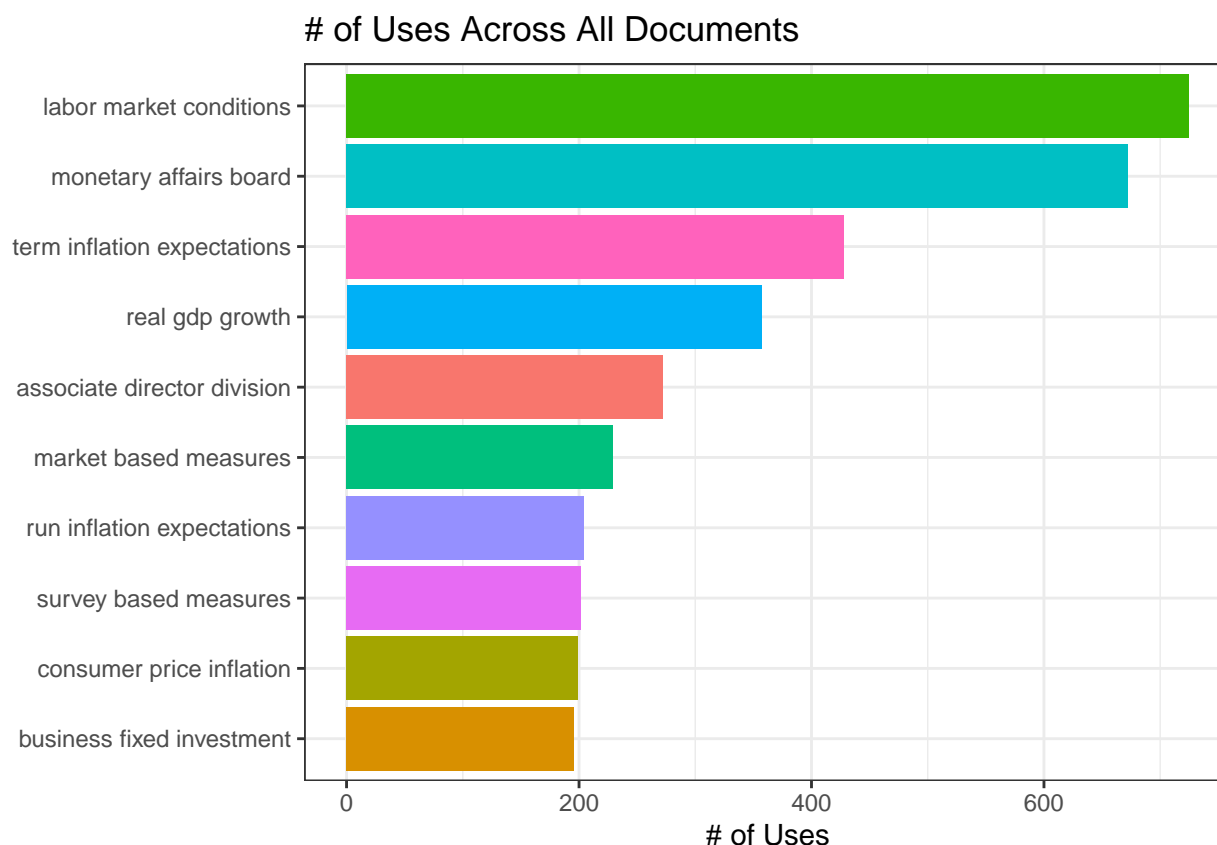


Using two words for the analysis allows us to ask more specific questions about the data that could not be answered with just one word. It allows us to ask things like, how has the Fed's use of "financial markets" compared with its use of "labor market" over time?

FOMC Mentions of "Financial Markets" vs. "Labor Market"



We can extend this line of thinking further with pairings of three words:



These terms might benefit from exploring combinations of four to find words that precede these phrases, such as “*strong* labor market conditions” or “*long* term inflation expectations.” This is not an optimal path to pursue, however, because the more combinations of terms there are to find, the fewer matches in the data there will be. Going from two to three terms made the “# of Uses” drop by about one-third. We can find these kinds of patterns another way. As we will soon see, networks find most of these combinations for us.

Sentiment Analysis


Now that we have terms stored as bigrams (and trigrams!), we can perform sentiment analysis on the Fed Minutes releases. Various lexicons have been built up over the years to determine the sentiments of a document. In this case, the Loughran and McDonald dictionary of financial sentiment words (Loughran and McDonald 2011) will be used. Other popular sentiment lexicons, such as the “AFINN,” “NRC,” or “BING” are not suited for documents coming from a bank—these three lexicons assign values to terms such as “share” or “risk” despite the fact those words may not carry any connotation in a financial setting.

Sentiment analysis could be used for the terms without sorting them into pairs, but this adds error when negation terms are used. For example, *weak* growth does not mean the same thing as *strong* growth, but this distinction would never be found by using only one term (the word “growth” does not appear in the loughran lexicon).

Whenever negation words preface a word in the loughran lexicon, we can flip its sign or set its sentiment value to zero, ensuring negation words do not skew the true sentiment values.

Before continuing to the sentiment analysis, it is important to note that sentiment lexicons do not contain the same number of positive words as negative words—there is an imbalance (in most cases, sentiment lexicons have more positive words than negative words—Loughran and McDonald’s lexicon is no different). The results,

then, should be considered in relation to one another rather than in an absolute sense. For comparisons to other variables, an index is likely more appropriate.



Look at that dip at the beginning of the 2020 recession! They don't call economics the dismal science for nothing...

Markov Chains

To visualize the relationship between the words further, we can arrange groups of words into a network.

Correlated Words

We can calculate Pearson correlation coefficients to determine how often particular words appear with one another. Examining correlation coefficients across each of the FOMC Meetings is not ideal since each one consists of thousands of words, many of them being used in each document. The most highly correlated words, broken down by document as opposed to by another tokenization, *that are not perfectly correlated* are still provided below, however. Many of the documents are more than ten pages long and refer to similar topics, so many of the words will have perfect correlation.

Table 2: Phi Coefficients for Whole Documents

item1	item2	correlation
participants	committee	-0.612
inflation	committee	-0.471
policy	participants	0.459
inflation	federal	-0.367
market	committee	0.343
policy	inflation	0.313
policy	committee	-0.281
participants	market	-0.210
participants	inflation	0.157
market	federal	0.140

As implied above, what should instead be done is break the documents into smaller pieces and calculate correlation coefficients for each of these subsections. To ensure proper nouns do not dominate the results with perfect correlation (and because R said my computer ran out of memory =/), only those terms which appear more than 100 times will be used to calculate the coefficients. Also note that a significant number of words were removed—words such as “secretary,” “jerome,” “tuesday,” “dallas,” “governors,” and many similar terms—to keep only the information we are interested in.

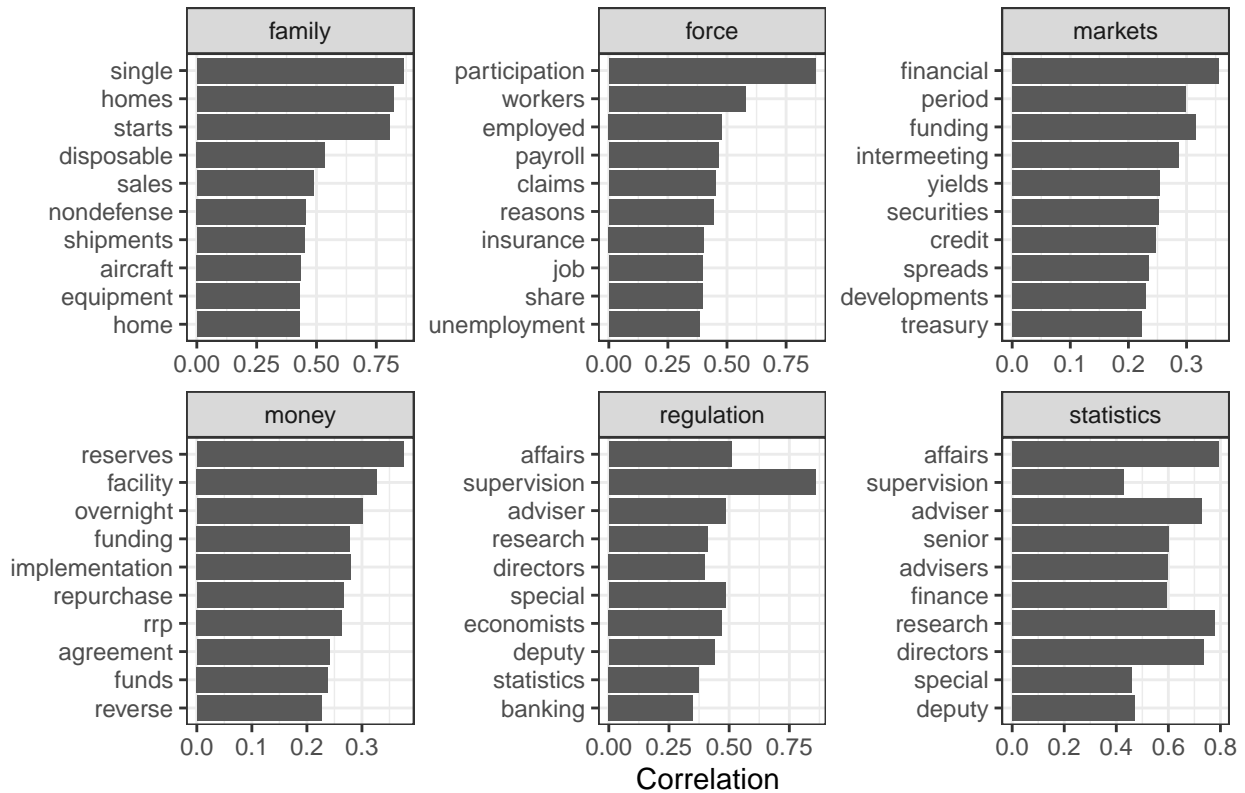
Table 3: Phi Coefficients for Document Sections

Word 1	Word 2	Correlation
speculative	grade	0.904
force	participation	0.875
notation	completed	0.871
single	family	0.864
supervision	regulation	0.859
accordance	authorize	0.832
shipments	nondefense	0.830
insurance	claims	0.830
homes	family	0.823
conclusion	authorize	0.811

The remainder of this section uses data from the above table.

Let’s see a nice graph showing the top 10 correlations of six words:

Most Correlated Terms for Each Word

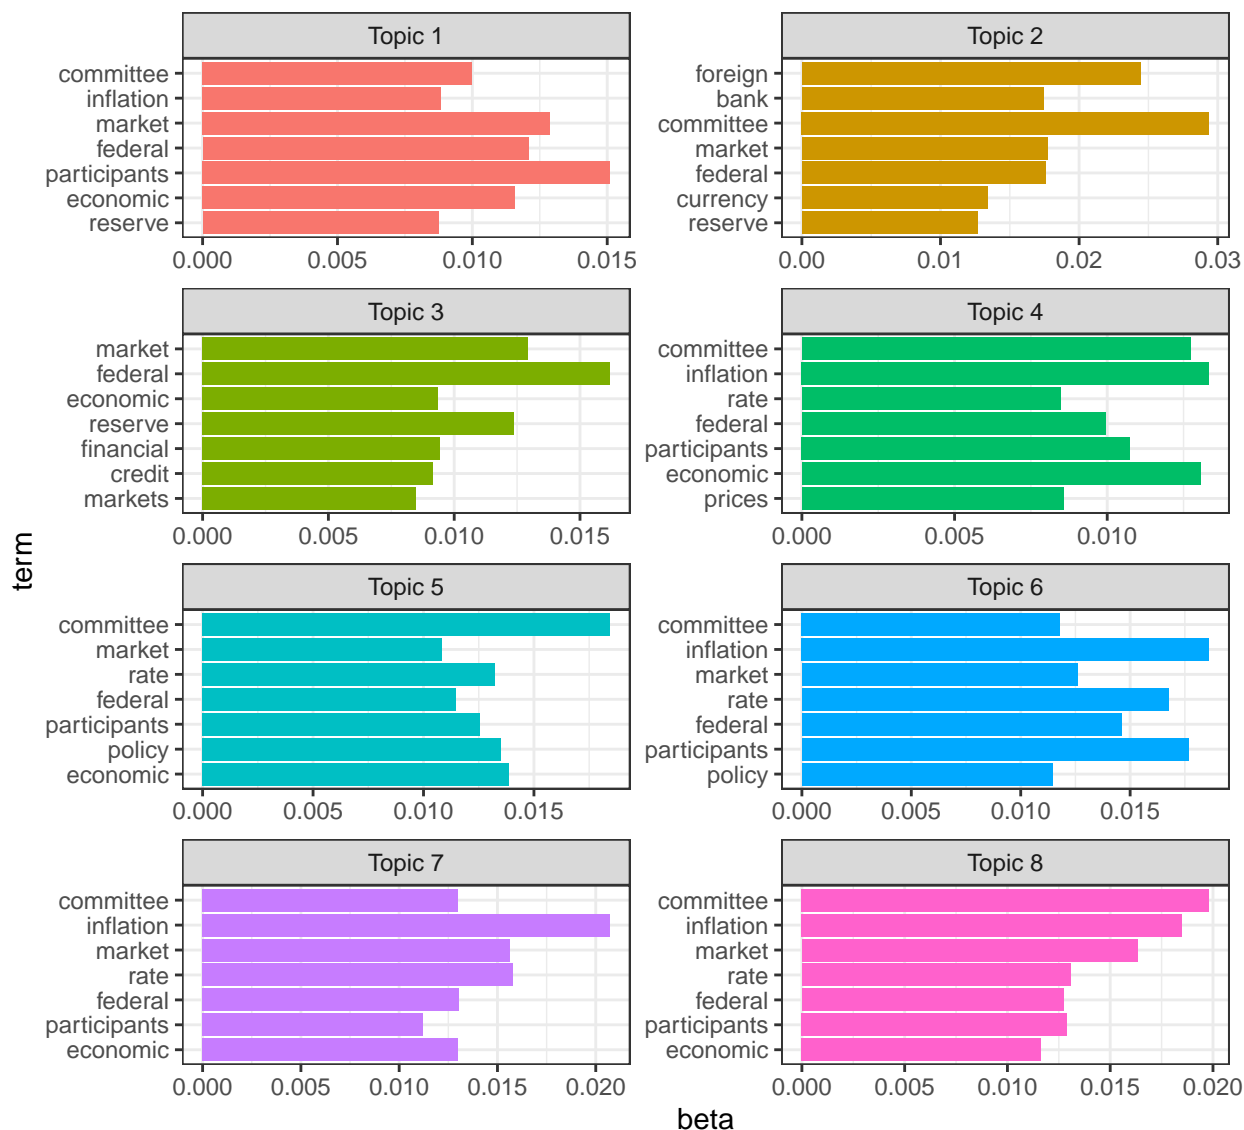


Pay attention to the scale and note that it is not the same for each of the six terms. These graphs paint a similar picture to the graphs from the N-grams section, though this one is a bit more telling because even though words that frequently appear next to each other are still included (Ex. “[labor] force participation”), words that often occur in the same sentence or set of sentences appear as well (Ex. “family” and “disposable” are unlikely to be written after one another in a document—which is the only way for n-grams to find a connection between the two—though correlation coefficients still recognize such co-occurrences).

One of the best uses of these correlations is to pick interesting words and find how other terms relate to them in a network, similar to that from the previous section. Note that because there is no directional relationship, there are no arrows on the graph. Since the number of words is enormous, only those terms with a correlation of at least 0.6 are included.


could be made up of terms such as “rate,” “yield,” etc. (*every topic is a mixture of words*).

You might have noticed that the word “rate” appeared in both the inflation and financial markets topics. That is fine—unlike other unsupervised learning/clustering methods, LDA allows words to be reused in different topics. Just as two words have different meanings in human language, two words can appear in different topics with LDA.



There is much overlap in the topics the model found. This is not surprising—the FOMC usually has a predictable agenda it needs to get through.


We may instead be interested in those words having the greatest difference between groups (as opposed to those most likely to appear in a particular topic). We can use a log ratio to calculate this difference. For example, which words are most different between topic 8 and topic 1? (For those unfamiliar with the math, just know that the negative values are terms belonging to topic 1 and the positive values belong to topic 8.)



If nothing else, we are all familiar with Topic 1. . .

In addition to finding which words belong to which topics, we can also determine which topics belong to which documents. Since we know roughly when Topic 1 is identified, we can confirm our observations and determine just when each topic is most likely to occur.

Probability Density Functions of Topics



Topic 2 appears in the first meeting of each year. There are 12 peaks in the sample period and 12 years of data. The meetings are not on January first, remember, and the vertical time bars on the charts have intervals of 2.5 years. There is a section at the beginning of the first FOMC Minutes release of each year called “Annual Organizational Matters” (see this FOMC Minutes release, for example) which only appears in first-meeting-of-the-year meetings. People are elected to their positions at the Fed (or are they appointed? I didn’t vote for them. . .). Topic 2 is the Fed’s version of housekeeping.

In the event you did not believe me, the probability of Topic 1 appearing in FOMC Minutes shoots up into the sky during the 2020 meetings.

References

Board of Governors of the Federal Reserve System (US), 10-Year Treasury Constant Maturity Rate [DGS10], retrieved from FRED, Federal Reserve Bank of St. Louis; <https://fred.stlouisfed.org/series/DGS10>, January 7, 2021.

Loughran, T. and McDonald, B. (2016). Textual analysis in accounting and finance: A survey. *Journal of Accounting Research* 54(4), 1187-1230. doi: 10.2139/ssrn.2504147 <https://sraf.nd.edu/textual-analysis/resources/#Master%20Dictionary>.